

Starling

Animated renderings of data structures for teaching

Manual for v0.2.0

Contents

1. Introduction	2
1.1. Quick start	2
2. Architecture	3
2.1. Snapshots, frames, and the renderer	4
2.2. Path identity	4
2.3. Sticky animation accumulation	4
2.4. The cetz integration	5
3. Design decisions	5
3.1. Why frames, not a wrap callback	5
3.2. Two-channel captions	6
3.3. The six-frame rotation	6
3.4. Path-based anchors and the cetz tree quirk	6
4. BST animations tour	7
4.1. Static display	7
4.2. Search	7
4.3. Insert	9
4.4. Delete	11
4.5. Rotate	12
4.6. Traversals	15
5. Touying composition examples	18
5.1. Captions in a sidebar	18
5.2. Driving layout from step metadata	19
6. Composing with cetz annotations	19
7. The Op command stream	20
8. API reference	21
8.1. Render helpers	21
8.2. Animation kernel	25

1. Introduction

Starling renders animated data structures in Typst, built on top of cetz, typsy, and (optionally) touying. The package is designed for a programming course: it ships a binary search tree today and is structured to grow to heaps, hash tables, and graphs without rewriting the animation kernel.


Every animation in starling is built from the same primitive — an ordered array of Frame records. Each Frame carries a rendered cetz canvas, an optional textual caption, and free-form step metadata. The package supplies a few helpers to collapse the array into a final image, a vertical stack, or an array of subslide figures, but you are free to ignore them and lay frames out however your document needs.

1.1. Quick start


```
#import "@preview/starling:0.2.0" as starling
#import starling: BST




#let t = (BST.new)(value: 4, label: auto, left: none, right: none)
#let t = (t.insert)(1)
#let t = (t.insert)(7)
#let t = (t.insert)(3)
#let t = (t.insert)(6)
```

A static render of the tree, taking the final frame from `(t.display)()`:



A full search animation, stacked vertically with per-step captions:





2. Architecture

Starling is layered into three source files, each with a single responsibility:

File	Role
src/tree-anim.typ	The animation kernel — defines Snapshot, Frame, TreeRenderer, draw-tree, Op, and apply-ops. Knows nothing about binary search trees specifically.

File	Role
src/bst.typ	The BST class — implements pure operations (insert, delete, rotate) and the <code>*-display</code> methods that build animations using the kernel.
src/lib.typ	The public surface — re-exports everything users need, plus the render helpers (last, stacked, figures, canvases-only).

This split lets us add a new data structure (e.g. a heap) by writing just a new file at the `bst.typ` layer; the kernel does not need to change. Conversely, a power user who wants to assemble custom animations can talk to the kernel directly without touching the BST class at all.

2.1. Snapshots, frames, and the renderer

Three types do most of the work. Their roles are intentionally distinct:

Type	Role
Snapshot	A <i>sparse style overlay</i> for one moment in time — which nodes are filled what colour, which edges are dashed or hidden, which notes are attached to which nodes. Built up by chaining <code>style-node</code> / <code>style-edge</code> / <code>note-node</code> calls.
TreeRenderer	A tree plus an ordered list of snapshots plus optional caption / step-metadata for each snapshot. The thing you build up while describing an animation.
Frame	The <i>output</i> record: (<code>canvas</code> , <code>caption</code> , <code>step</code>). <code>TreeRenderer.render</code> produces an array of these, one per snapshot. This is what <code>*-display</code> methods return.

The data flow is:

Snapshot array → `TreeRenderer.render` → `Array(Frame)` → helper → document content

2.2. Path identity

Every node is identified by its position in the tree, encoded as a string of "L" and "R" characters from the root. The root is "", the right child of the root is "R", the left child of "R" is "RL", and so on. Edges are identified by the path of their *child* node — each child has exactly one parent, so this is unambiguous.

Path identity is the linchpin that keeps node and edge styling independent of the tree's value or layout. It also means a node-highlighting animation built for one tree can be reused on a tree of the same shape with different values just by passing a different tree to `make-renderer`.

The only places that interpret path characters are `PathId` (the refined string type), `_build-cetz-tree` (which walks `.left` and `.right`), and the BST's own `by-value` / `path-to` / `resolve` methods. Everywhere else treats paths as opaque keys, which leaves room to generalise — n-ary trees would use slash-separated child indices like "0/1/2" or arrays of ints, and only those four places need to change.

2.3. Sticky animation accumulation

`make-renderer(tree, sticky: true)` is the common case: each new frame pushed via `r.push-frame()` starts from the previous frame's snapshot. That means a sequence like

```
#let r = starling.make-renderer(t, sticky: true)
#let r = (r.push-with-node)("", stroke: blue + 2pt) // frame 1
```

```
#let r = (r.push-with-node)("L", stroke: blue + 2pt) // frame 2
#let r = (r.push-with-node)("LR", stroke: blue + 2pt) // frame 3
```

produces three frames where each highlights one more node than the last — frame 3 has all three highlights, not just LR. This matches the natural mental model for teaching (“we visited these nodes in order”). Set `sticky: false` if you want each frame to start fresh.

Captions and step metadata do *not* sticky-accumulate; each frame’s caption is whatever was last set on that frame, with no carry-over.

2.4. The `cetz` integration

The animation kernel renders each snapshot via `draw-tree`, which emits `cetz` draw commands using `cetz.tree.tree(...)` for layout. The key choice here is *not* to wrap `draw-tree`’s output in `cetz.canvas` inside the kernel — that wrap happens one layer up, in `_render-canvas`. The two functions are split so a power user can write their own `cetz.canvas` block, call `draw-tree` inside it, and add their own `cetz` annotations alongside the tree (see Section 6).

Anchors for individual nodes are exposed by `path-anchor`, which translates an L/R path string into the `cetz` anchor name produced by `draw-tree`. The fact that this needs a translation function at all is a quirk worth flagging — see Section 3.4.

3. Design decisions

This section explains the major calls that shape the API. None of them are obvious from reading the code, but each was a conscious choice and the reasoning is load-bearing for anyone considering a refactor.

3.1. Why frames, not a wrap callback

The original (v0.1) API let users configure `starling.configure(wrap: alternatives)` to bake `touying`’s `alternatives` into the `BST` class via closure capture. Each `*-display` method internally called `wrap(..figs)` on the array of figures it generated.

That approach was rigid in two ways. First, the rendering decision was made *inside* the animation method, so the same animation could not appear as a static figure in handouts and as a subslide sequence in slides — you’d configure differently per context. Second, and worse, the animation was always one opaque blob: there was no way to weave the caption “ $6 > 4$ ” alongside other slide content while keeping the visual animation synchronised, because everything lived inside a single `alternatives` mark.

The current API solves both: `*-display` returns the raw frame array, and the helpers (`last`, `stacked`, `figures`) collapse it however the caller wants. Wanting to lay captions in a sidebar while the animation plays? Pull them out yourself:

```
#let frames = (t.search-display)(6)
#grid(columns: 2,
  alternatives(..starling.figures(frames, caption: false)),
  alternatives(..frames.map(f => f.caption)),
)
```

A historical note worth preserving: closure-captured `wrap` was itself chosen over a state-based approach (`std.state`, `typsy`’s `safe-state`) because `touying`’s `alternatives` is a layout-time “mark” that `touying`

rejects inside context `{ ... }` blocks, and reading from state requires being inside a context. With the new frames-based API this constraint goes away entirely — frames are plain data, no state, no context.

3.2. Two-channel captions

For search-display and insert-display, the comparison string (" $6 > 4$ ") appears in *two* places:

- **Inline**, as a small note drawn beside the highlighted node in the cetz canvas. This labels which comparison happened at which node spatially.
- **Caption**, as a textual track on the Frame record. This is what stacked and figures render below the canvas, and what callers pulling captions out for custom layouts read.

The redundancy is deliberate. The inline note is part of the visual: when looking at a single frame, you should be able to see at a glance what the algorithm just did. The caption is a parallel *textual* track for layouts that want narration apart from the animation.

This is why `last(frames)` defaults to `caption: false` — the inline note already labels the node, so adding a caption block below the static figure would be redundant. `stacked` and `figures` default to `caption: true` because animated rendering invites the textual narrative.

3.3. The six-frame rotation

`rotate-display` walks through rotation as six discrete frames: `init`, `pivots`, `break`, `restructure`, `connect`, `settle`. The sequence is calibrated for teaching: each frame answers exactly one question.

step.kind	Question answered
init	What's the starting tree?
pivots	Which two nodes are rotating?
break	Which edges are about to disappear?
restructure	What does the new shape look like, edges aside?
connect	Where do the new edges go?
settle	What does the final, clean tree look like?

An earlier draft had a “dashed red” intermediate between `pivots` and `break` (edges going dashed before disappearing). That was dropped because the orange pivot highlight already cues the eye to those edges — the extra frame added length without information.

The `restructure` frame is the load-bearing one: it shows the new tree shape with the rotated edges still hidden, so the student can see the new positions form *before* the edges connect. Without it, the rotation would collapse into one cut from “broken” to “rotated + green edges”, which is too much information at once.

3.4. Path-based anchors and the cetz tree quirk

`cetz.tree` numbers its nodes internally as `0`, `0-0`, `0-1`, `0-0-1`, ... — sibling indices joined by hyphens, with the root being `0`. Starling configures `cetz-tree`'s `group-name-prefix` to `"node-"`, so the resulting anchor names look like `node-0-0-1`. The whole tree also sits inside an outer named group (`"tree"` by default), so the fully qualified anchor for path `"LR"` is `tree.node-0-0-1`.

Two things made me reach for path-anchor rather than letting users type that out:

1. The user thinks in "LR", not "0-0-1". The translation is mechanical ($L \rightarrow 0, R \rightarrow 1$, prepend the root 0) but it's error-prone to do by hand.
2. The naming scheme is a `cetz-tree` implementation detail. Wrapping it in `path-anchor(path, tree-name:, prefix:)` lets us swap to a different naming scheme later (e.g. if we drop `cetz-tree` for a custom layout) without churn at every annotation call site.

Phantom siblings (used to off-centre lone children — see `_build-cetz-tree`) also get `cetz-tree-internal` anchor names like `node-0-0-0`, but they are hidden at draw time and you generally should not annotate them.

4. BST animations tour


The `BST` class provides five `*-display` methods. Each returns `Array(Frame)`; the examples below pass the result to `starling.stacked` to render every frame vertically with its caption.

We'll use the same sample tree throughout this section:

```
#let t = (BST.new)(value: 4, label: auto, left: none, right: none)
#let t = (t.insert-many)(1, 7, 3, 6, 8)
```




4.1. Static display


`(t.display)()` returns a one-element frame array — the unmodified tree, with no step metadata.



4.2. Search


`(t.search-display)(v)` walks the search path, highlighting each node visited and labelling it with the comparison made. `step.kind` is "init" for the initial frame and "compare" for each step; `step.found` records whether the value was reached.








4.3. Insert

`(t.insert-display)(v)` walks the search path as for `search-display`, then transitions to the after-tree with the new node highlighted green. `step.kind` is "init", "compare" (per search step), and finally "inserted".








4.4. Delete

`(t.delete-display)(v)` has three internal cases — leaf, one-child, and two-children — and dispatches on the target's children. `step.kind` ranges over "init", "highlight", "break", "descend", "transfer", and "settle" depending on case.


The two-children deletion is the most complex: it highlights the target, descends into the left subtree to find the in-order predecessor, annotates the value transfer, then jumps to the after-tree with the new root of the affected subtree highlighted green.




Delete 4




Find predecessor



Find predecessor





Transfer 3




Done

4.5. Rotate


`(t.rotate-display)(c)` rotates around the direct child whose value is `c`. The six-frame sequence is described in Section 3.3.




Rotate around 1




Break edges



Restructure tree




Reconnect edges









4.6. Traversals

`(t.in-order-display)()`, `(t.pre-order-display)()`, `(t.post-order-display)()`, and `(t.level-order-display)()` animate the four standard traversals. Each returns one frame per visit (plus the initial frame): the visited node is filled with the next color in a perceptually-uniform palette (magma by default) and tagged with a numbered badge marking its visit order, and the caption accumulates the running output sequence. `step.kind` is "init" for the initial frame and "visit" thereafter, with `step.value` and `step.index` (1-indexed) recording each visit. Pass a different palette: argument (e.g. `palette: color.map.viridis`) to switch color maps.


Pure-data variants `(t.in-order)()`, `(t.pre-order)()`, `(t.post-order)()`, and `(t.level-order)()` return the visit sequence as an array of path strings without producing any animation, in case you want to drive a custom layout.







The final frame of each traversal carries the algorithm's full color signature. Laying all four out side-by-side gives a fast visual comparison — especially useful as a recall cue in review material where students already understand the algorithms and the gradient pattern reads as “oh right, post-order goes leaves-cool to root-warm”:



5. Touying composition examples

Starling is layout-agnostic. In a touying deck, the simplest use splats figures into alternatives:

```
#import "@preview/touying:0.7.3": *
#import "@preview/starling:0.2.0" as starling
#import starling: BST

== Searching
#alternatives(..starling.figures((t.search-display)(6)))
```

Each frame becomes a subslide; touying handles the rest.

5.1. Captions in a sidebar

When the visual animation should sit alongside synchronised text elsewhere on the slide, pull the captions out separately:

== Searching for 6

```
#let frames = (t.search-display)(6)
#grid(columns: 2, column-gutter: 2em,
  alternatives(..starling.figures(frames, caption: false)),
  alternatives(..frames.map(f => f.caption)),
)
```

Both alternatives calls produce the same number of subslides, so they step in lockstep.

5.2. Driving layout from step metadata

frame.step is free-form per-method metadata. For example, a slide that wants to colour-code its narration by step kind:

```
#let kind-colors = (compare: blue, inserted: green)

#let captioned-frames = frames.map(f => {
  let color = kind-colors.at(f.step.kind, default: black)
  figure(
    stack(dir: ttb, spacing: 0.5em,
      f.canvas, text(fill: color, f.caption)),
  )
})


#alternatives(..captioned-frames)
```

6. Composing with cetz annotations

For callouts or overlays that need to track specific nodes, drop down to the cetz layer. draw-tree(tree, snapshot, ...) emits the tree's cetz drawables without wrapping them in cetz.canvas, so you can compose them with your own draw commands in a shared canvas. path-anchor(path) translates an L/R path string to the cetz anchor name of the corresponding node.

```
#import "@preview/cetz:0.5.2"

#cetz.canvas({
  starling.draw-tree(t, starling.blank-snapshot())
  import cetz.draw: *
  circle(starling.path-anchor("LR"), radius: 0.85, stroke: red + 2pt)
})
```





The same trick combines with frame canvases: each frame’s canvas is itself a `cetz.canvas` block, so for static “annotate the final frame” use you can either render once via `draw-tree` (as above) or extract the `cetz` block from a frame and add to it externally.

7. The Op command stream

For animations that don’t fit the BST’s built-in methods — say, a custom hash-table probe sequence or a graph traversal — drop down to the `Op` command stream. Build a sequence of declarative ops and fold them into a renderer with `apply-ops`:

```

#let r = starling.make-renderer(t, sticky: true)
#let r = (r.with-caption)([start at root])
#let r = starling.apply-ops(r, (
  (starling.Op.Highlight.new)(path: "", color: blue),
  (starling.Op.Annotate.new)(path: "", text: [1 < 4]),
  (starling.Op.Commit.new)(
    alt: "Comparing the target 1 against the root 4.",
  ),
),
)
#let r = (r.with-caption)([descend left])
#let r = starling.apply-ops(r, (
  (starling.Op.ClearNotes.new)(),
  (starling.Op.Highlight.new)(path: "L", color: blue),
  (starling.Op.StyleEdge.new)(path: "L", style: (stroke: blue + 2pt)),
  (starling.Op.StyleNode.new)(path: "L", style: (fill: green.lighten(70%))),
  (starling.Op.Annotate.new)(path: "L", text: [1 = 1]),
  (starling.Op.Alt.new)(
    text: "Descended into the left subtree; 1 matches the left child.",
  ),
),
)
#starling.stacked((r.render)())
  
```



`Op.Commit` closes the current frame, attaches the supplied `alt` text to it, and opens a fresh blank frame. The `alt` argument is required so every frame the command stream produces carries accessible text. The trailing in-progress frame is kept implicitly — no final `Op.Commit` is needed — but use `Op.Alt(text)` (or `r.with-alt(...)` on the returned renderer) to give that last frame its `alt` text too.

`Op.Caption` does *not* exist; captions and step metadata are renderer-level, set via `r.with-caption(...)` / `r.with-step(...)` directly between `Op` batches. The `split` keeps each `Op`'s responsibility narrow (modifying the current snapshot) and matches how the `*-display` methods are written internally.

8. API reference

The remainder of this manual is an auto-generated reference, produced by the `tidy` package from doc-comments in the source.

8.1. Render helpers

- `canvases-only()`
- `figures()`
- `last()`
- `stacked()`

Variables

- `BST`

8.1.1. canvases-only

Drop captions and step metadata, returning just the array of canvases. Useful when you want to lay out captions yourself — e.g. alongside other slide content in a custom touying layout.

8.1.1.1. Parameters

```
canvases-only(frames: array) -> array
```

8.1.1.1.1. frames array

The frame array returned by any `*-display` method.

8.1.2. figures

Build an array of figure content, one per frame, suitable for splatting into touying's `alternatives(...)` so each frame becomes its own subslide. Starling itself does not depend on touying — the result is just an array of standard Typst figures.

8.1.2.1. Parameters

```
figures(  
  frames: array,  
  caption: bool,  
  alt: auto str,  
  caption-spacing: length  
) -> array
```

8.1.2.1.1. frames array

The frame array returned by any `*-display` method.

8.1.2.1.2. caption bool

Whether to stack each frame's caption below its canvas inside the figure body.

Default: `true`

8.1.2.1.3. alt auto or str

Alt-text override applied to every frame. `auto` uses each frame's `alt` field, falling back to the caption (if it's a plain string) and then to a generic "Animation frame N" placeholder.

Default: `auto`

8.1.2.1.4. caption-spacing `length`

Spacing between canvas and caption when `caption: true`.

Default: `0.5em`

8.1.3. last

Final frame as a single piece of content. Useful for static renderings in print or for “just show me the end state” usage.

The inline node annotations (drawn next to nodes in the `cetz` canvas) already label what the final step did, so the textual caption is suppressed by default. Pass `caption: true` to stack the caption below the canvas.

The returned content is wrapped in an alt-tagged figure for accessibility. The alt text comes from the frame’s `alt` field; pass an explicit string to `alt` to override.

8.1.3.1. Parameters

```
last(  
  frames: array,  
  caption: bool,  
  spacing: length,  
  alt: auto str  
) -> content
```

8.1.3.1.1. frames `array`

The frame array returned by any `*-display` method.

8.1.3.1.2. caption `bool`

Whether to render the step’s caption below the canvas.

Default: `false`

8.1.3.1.3. spacing `length`

Vertical spacing between canvas and caption when `caption: true`.

Default: `0.5em`

8.1.3.1.4. alt auto or str

Alt-text override. auto uses the frame's alt field.

Default: auto

8.1.4. stacked

All frames stacked vertically as one block of content. Useful for handouts that want to show the full animation in a single figure. Captions are on by default since the stack is read as a sequence — the caption text annotates which step each tree corresponds to.

Each frame is wrapped individually in an alt-tagged figure so screen-reader users hear the per-step narration. Pass a string to alt to set the same override on every frame.

8.1.4.1. Parameters

```
stacked(  
  frames: array,  
  caption: bool,  
  spacing: length,  
  caption-spacing: length,  
  alt: auto str  
) -> content
```

8.1.4.1.1. frames array

The frame array returned by any *-display method.

8.1.4.1.2. caption bool

Whether to include each frame's caption below its canvas.

Default: true

8.1.4.1.3. spacing length

Vertical spacing between frames.

Default: 1em

8.1.4.1.4. caption-spacing length

Spacing between each canvas and its caption.

Default: 0.5em

8.1.4.1.5. alt auto or str

Alt-text override applied to every frame. auto uses each frame's alt field.

Default: auto

8.1.5. BST

The default BST class. Build a tree via (BST.new) (value:, label:, left:, right:) — set label: auto to render the integer value as the drawn label, or pass content (a string, image, or arbitrary content) for custom labels. Call its *-display methods (see the BST API section of the manual) to get animation frames; pass those frames through one of the helpers below.

8.2. Animation kernel

- apply-ops()
- blank-snapshot()
- concat-frames()
- draw-tree()
- make-renderer()
- path-anchor()

Variables

- PathId
- NodeStyle
- EdgeStyle
- Snapshot
- Frame
- TreeRenderer
- Op

8.2.1. apply-ops

Fold a sequence of Op values into a renderer, returning the updated renderer. The trailing in-progress frame is kept — no explicit Op.Commit is needed after the last batch of edits, but attach alt text to that trailing frame via Op.Alt (or r.with-alt(...) on the returned renderer) so it has accessible text like the committed frames do.

8.2.1.1. Parameters

```
apply-ops(  
  renderer: TreeRenderer,  
  ops: array  
) -> TreeRenderer
```

8.2.1.1.1. renderer TreeRenderer

The renderer to apply ops to.

8.2.1.1.2. `ops` `array`

Sequence of Op values to fold left-to-right.

8.2.2. `blank-snapshot`

Build a fresh Snapshot with no node or edge style overrides. Useful when you want to render a tree once with `draw-tree` and no per-frame state.

8.2.2.1. Parameters

`blank-snapshot()` -> `Snapshot`

8.2.3. `concat-frames`

Stitch frame arrays from multiple renderers (or pre-rendered arrays) into one flat `Array(Frame)`. Use this when an animation spans a tree-shape change (rotation, deletion, insertion) and one renderer can't hold all the frames — one renderer per shape, joined at the boundary.

Accepts a mix of `TreeRenderer` instances and already-rendered arrays; the former are render-ed in place.

8.2.3.1. Parameters

`concat-frames(..parts)` -> `array`

8.2.3.1.1. `..parts`

Variadic mix of renderers and frame arrays.

8.2.4. `draw-tree`

Emit the `cetz` draw commands for one styled snapshot of `tree`, *without* wrapping them in a `cetz.canvas`. Caller is responsible for the canvas wrap. Use this when you want to add your own `cetz` annotations alongside the tree — for example, callouts anchored at specific nodes.

The whole tree is wrapped in a named `cetz` group (name, default `"tree"`); each non-phantom node has an inner-name of `<node-prefix><cetz-tree-name>` where the `cetz-tree` name is `"0"` for the root, `"0-0"` for L, `"0-1"` for R, `"0-0-1"` for LR, and so on. Fully qualified anchor names therefore look like `"tree.node-0-0-1"` — use `path-anchor()` to compute them from an L/R path string.

8.2.4.1. Parameters

```
draw-tree(  
    tree: dictionary,  
    snapshot: Snapshot,  
    default-node-style: dictionary,  
    default-edge-style: dictionary,  
    name: str,  
    node-prefix: str  
) -> content
```

8.2.4.1.1. tree dictionary

The tree to render — any value with value, label, left, and right fields (e.g. a BST instance). label of auto falls back to str(value) any other value (string, image, content) is rendered as the node's drawn label.

8.2.4.1.2. snapshot Snapshot

Style overlay for this snapshot. Use blank-snapshot() for an unstyled tree.

8.2.4.1.3. default-node-style dictionary

Default node-style overrides applied before per-node overrides.

Default: (:)

8.2.4.1.4. default-edge-style dictionary

Default edge-style overrides applied before per-edge overrides.

Default: (:)

8.2.4.1.5. name str

Outer group name for the whole tree. Must match the tree-name argument to path-anchor() for anchors to resolve.

Default: "tree"

8.2.4.1.6. node-prefix str

Prefix attached to each node's cetz-tree internal name. Must match the prefix argument to path-anchor().

Default: "node- "

8.2.5. make-renderer

Build a `TreeRenderer` seeded with one blank initial frame.

With `sticky: true` (the default), each new frame pushed via `r.push-frame()` starts from the previous frame's style snapshot — so highlights accumulate over the course of an animation. Set `sticky: false` to start each frame from a clean slate.

8.2.5.1. Parameters

```
make-renderer(  
  tree: dictionary,  
  default-node-style: dictionary,  
  default-edge-style: dictionary,  
  sticky: bool  
) -> TreeRenderer
```

8.2.5.1.1. tree dictionary

The tree to render — any value with `value`, `label`, `left`, and `right` fields. See `draw-tree()` for how `label` is rendered.

8.2.5.1.2. default-node-style dictionary

Default node-style overrides applied before per-snapshot overrides.

Default: `(:)`

8.2.5.1.3. default-edge-style dictionary

Default edge-style overrides applied before per-snapshot overrides.

Default: `(:)`

8.2.5.1.4. sticky bool

When `true`, new frames inherit the previous frame's style. When `false`, each new frame starts blank.

Default: `true`

8.2.6. path-anchor

Translates a starling path-id (L/R string rooted at `"`) to the fully-qualified cetz anchor name produced by `draw-tree()`. Path `"` maps to the root anchor; `"L"/"R"` map to first-level children; and so on.

The `tree-name` and `prefix` arguments must match the name and node-prefix passed to `draw-tree`.

8.2.6.1. Parameters

```
path-anchor(  
  path: str,  
  tree-name: str,  
  prefix: str  
) -> str
```

8.2.6.1.1. path str

L/R path string identifying a node (" " = root, "LR" = root's left child's right child, etc.).

8.2.6.1.2. tree-name str

Outer-group name; must match draw-tree's name.

Default: "tree"

8.2.6.1.3. prefix str

Per-node prefix; must match draw-tree's node-prefix.

Default: "node- "

8.2.7. PathId

Typsy refinement: a string built only from "L" and "R" characters. The empty string is the root; "L" is the root's left child; "LR" is the root's left child's right child; and so on.

8.2.8. NodeStyle

Typsy refinement: a dictionary of node-style overrides. Recognised keys are fill, stroke, text-fill, note, note-fill, and hide.

8.2.9. EdgeStyle

Typsy refinement: a dictionary of edge-style overrides. Recognised keys are stroke, note, note-fill, mark, and hide.

8.2.10. Snapshot

Typsy class representing one sparse style overlay — the per-node and per-edge style overrides for a single frame. Build one via `blank-snapshot()` and chain its `style-node` / `style-edge` / `note-node` / `note-edge` / `clear-notes` methods. Snapshots are rarely constructed directly — they’re managed by `make-renderer()` and the BST’s `*-display` methods.

8.2.11. Frame

Typsy class representing one rendered animation frame. Fields: `canvas` (cetz content, no caption baked in), `caption` (optional textual track for the step, possibly none), `step` (optional free-form per-method metadata, possibly none), and `alt` (optional accessible text describing what the frame depicts, possibly none). Produced by `TreeRenderer.render()` the BST’s `*-display` methods return arrays of these.

The `alt` field carries text destined for the `alt:` argument on a Typst figure. The first frame of each operation includes the full tree structure (via the BST `describe` method) so screen-reader users have a baseline; subsequent frames describe only the per-step change.

8.2.12. TreeRenderer

Typsy class accumulating an animation. Holds the tree, parallel arrays of snapshots / captions / steps / alts (one entry per frame), and default node/edge styles. Methods include `push-frame()`, `patch(fn)`, `push-with-node(path, ..style)`, `push-with-edge(path, ..style)`, `push-note-node(path, txt)`, `push-note-edge(path, txt)`, `with-caption(c)`, `with-step(s)`, `with-alt(a)`, and `render()`. Build one via `make-renderer()`.

8.2.13. Op

Typsy enumeration describing a declarative command stream for building animations. Variants:

- `Op.Highlight(path, color)` — set a node’s stroke to `color + 2pt`.
- `Op.Annotate(path, text)` — attach an inline note to a node (drawn in-canvas).
- `Op.StyleNode(path, style)` — arbitrary node-style overrides.
- `Op.StyleEdge(path, style)` — arbitrary edge-style overrides.
- `Op.Commit(alt)` — finalise the current frame, attach `alt` text to it, and open a fresh blank frame. `alt` is required so every frame the command stream produces carries accessible text.
- `Op.Alt(text)` — set `alt` text on the in-progress frame without committing. Primary use is the trailing frame, which has no following `Op.Commit` to carry its `alt`; can also be used to set `alt` earlier in a frame’s lifecycle if convenient.
- `Op.ClearNotes()` — drop all notes from the current frame’s snapshot.

Captions and step metadata are renderer-level rather than snapshot-level: set them between `Op` batches via `r.with-caption(...)` / `r.with-step(...)`, not through an `Op`. Fold a sequence of `Ops` into a renderer with `apply-ops()`.