Starling

Animated renderings of data structures for teaching

Manual for v0.2.0

Contents

1. Introduction

Starling renders animated data structures in Typst, built on top of cetz, typsy, and (optionally) touying. The package is designed for a programming course: it ships a binary search tree today and is structured to grow to heaps, hash tables, and graphs without rewriting the animation kernel.

Every animation in starling is built from the same primitive — an ordered array of `Frame` records. Each `Frame` carries a rendered cetz canvas, an optional textual caption, and free-form step metadata. The package supplies a few helpers to collapse the array into a final image, a vertical stack, or an array of subslide figures, but you are free to ignore them and lay frames out however your document needs.

1.1. Quick start

```
#import "@preview/starling:0.2.0" as starling
#import starling: BST

#let t = (BST.new)(value: 4, label: auto, left: none, right: none)
#let t = (t.insert)(1)
#let t = (t.insert)(7)
#let t = (t.insert)(3)
#let t = (t.insert)(6)
```

A static render of the tree, taking the final frame from `(t.display)()`:

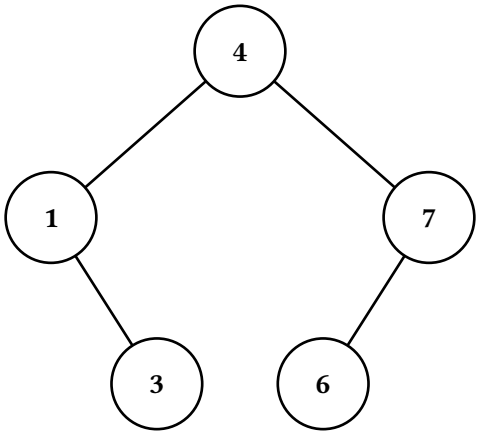

A full search animation, stacked vertically with per-step captions:

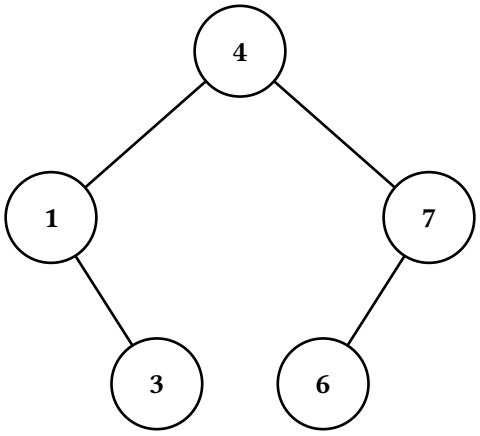

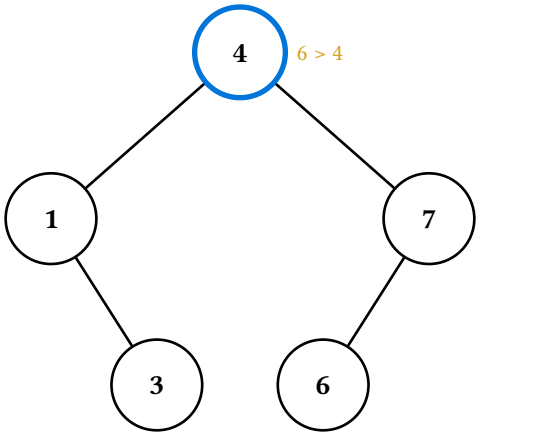

6 > 4

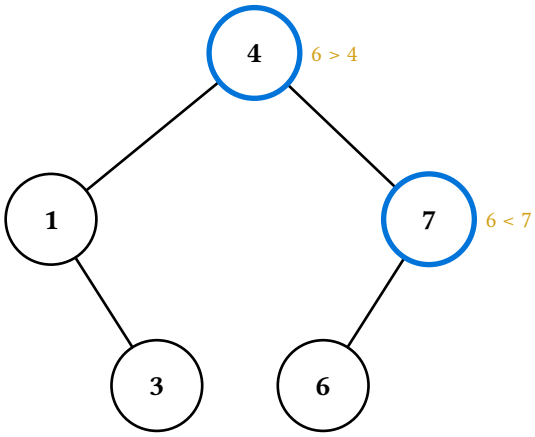

6 < 7

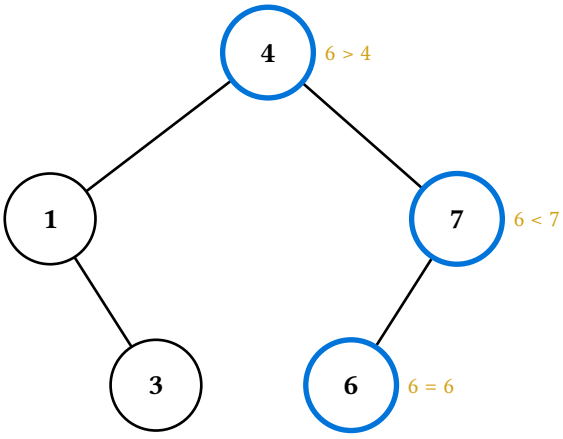

6 = 6

2. Architecture

Starling is layered into three source files, each with a single responsibility:

File	Role
`src/tree-anim.typ`	The animation kernel — defines `Snapshot`, `Frame`, `TreeRenderer`, `draw-tree`, `Op`, and `apply-ops`. Knows nothing about binary search trees specifically.

File	Role
`src/bst.typ`	The BST class — implements pure operations (`insert`, `delete`, `rotate`) and the `*-display` methods that build animations using the kernel.
`src/lib.typ`	The public surface — re-exports everything users need, plus the render helpers (`last`, `stacked`, `figures`, `canvases-only`).

This split lets us add a new data structure (e.g. a heap) by writing just a new file at the `bst.typ` layer; the kernel does not need to change. Conversely, a power user who wants to assemble custom animations can talk to the kernel directly without touching the BST class at all.

2.1. Snapshots, frames, and the renderer

Three types do most of the work. Their roles are intentionally distinct:

Type	Role
`Snapshot`	A *sparse style overlay* for one moment in time — which nodes are filled what colour, which edges are dashed or hidden, which notes are attached to which nodes. Built up by chaining `style-node` / `style-edge` / `note-node` calls.
`TreeRenderer`	A tree plus an ordered list of snapshots plus optional caption / step-metadata for each snapshot. The thing you build up while describing an animation.
`Frame`	The *output* record: `(canvas, caption, step)`. `TreeRenderer.render` produces an array of these, one per snapshot. This is what `*-display` methods return.

The data flow is:

`Snapshot` array ⟶ `TreeRenderer.render` ⟶ `Array(Frame)` ⟶ helper ⟶ document content

2.2. Path identity

Every node is identified by its position in the tree, encoded as a string of `"L"` and `"R"` characters from the root. The root is `""`, the right child of the root is `"R"`, the left child of `"R"` is `"RL"`, and so on. Edges are identified by the path of their *child* node — each child has exactly one parent, so this is unambiguous.

Path identity is the linchpin that keeps node and edge styling independent of the tree's value or layout. It also means a node-highlighting animation built for one tree can be reused on a tree of the same shape with different values just by passing a different `tree` to `make-renderer`.

The only places that interpret path characters are `PathId` (the refined string type), `_build-cetz-tree` (which walks `.left` and `.right`), and the BST's own `by-value` / `path-to` / `resolve` methods. Everywhere else treats paths as opaque keys, which leaves room to generalise — n-ary trees would use slash-separated child indices like `"0/1/2"` or arrays of ints, and only those four places need to change.

2.3. Sticky animation accumulation

`make-renderer(tree, sticky: true)` is the common case: each new frame pushed via `r.push-frame()` starts from the previous frame's snapshot. That means a sequence like

```
#let r = starling.make-renderer(t, sticky: true)
#let r = (r.push-with-node)("",  stroke: blue + 2pt)   // frame 1
```

```
#let r = (r.push-with-node)("L", stroke: blue + 2pt)   // frame 2
#let r = (r.push-with-node)("LR", stroke: blue + 2pt)  // frame 3
```

produces three frames where each highlights one more node than the last — frame 3 has all three highlights, not just `LR`. This matches the natural mental model for teaching ("we visited these nodes in order"). Set `sticky: false` if you want each frame to start fresh.

Captions and `step` metadata do *not* sticky-accumulate; each frame's caption is whatever was last set on that frame, with no carry-over.

2.4. The cetz integration

The animation kernel renders each snapshot via `draw-tree`, which emits cetz draw commands using `cetz.tree.tree(...)` for layout. The key choice here is *not* to wrap `draw-tree`'s output in `cetz.canvas` inside the kernel — that wrap happens one layer up, in `_render-canvas`. The two functions are split so a power user can write their own `cetz.canvas` block, call `draw-tree` inside it, and add their own cetz annotations alongside the tree (see Section 6).

Anchors for individual nodes are exposed by `path-anchor`, which translates an L/R path string into the cetz anchor name produced by `draw-tree`. The fact that this needs a translation function at all is a quirk worth flagging — see Section 3.4.

3. Design decisions

This section explains the major calls that shape the API. None of them are obvious from reading the code, but each was a conscious choice and the reasoning is load-bearing for anyone considering a refactor.

3.1. Why frames, not a wrap callback

The original (v0.1) API let users configure `starling.configure(wrap: alternatives)` to bake touying's `alternatives` into the BST class via closure capture. Each `*-display` method internally called `wrap(..figs)` on the array of figures it generated.

That approach was rigid in two ways. First, the rendering decision was made *inside* the animation method, so the same animation could not appear as a static figure in handouts and as a subslide sequence in slides — you'd configure differently per context. Second, and worse, the animation was always one opaque blob: there was no way to weave the caption "6 > 4" alongside other slide content while keeping the visual animation synchronised, because everything lived inside a single `alternatives` mark.

The current API solves both: `*-display` returns the raw frame array, and the helpers (`last`, `stacked`, `figures`) collapse it however the caller wants. Wanting to lay captions in a sidebar while the animation plays? Pull them out yourself:

```
#let frames = (t.search-display)(6)
#grid(columns: 2,
  alternatives(..starling.figures(frames, caption: false)),
  alternatives(..frames.map(f => f.caption)),
)
```

A historical note worth preserving: closure-captured `wrap` was itself chosen over a state-based approach (`std.state`, typsy's `safe-state`) because touying's `alternatives` is a layout-time "mark" that touying

rejects inside `context { ... }` blocks, and reading from state requires being inside a `context`. With the new frames-based API this constraint goes away entirely — frames are plain data, no state, no `context`.

3.2. Two-channel captions

For `search-display` and `insert-display`, the comparison string ("6 > 4") appears in *two* places:

- **Inline**, as a small note drawn beside the highlighted node in the cetz canvas. This labels which comparison happened at which node spatially.
- **Caption**, as a textual track on the `Frame` record. This is what `stacked` and `figures` render below the canvas, and what callers pulling captions out for custom layouts read.

The redundancy is deliberate. The inline note is part of the visual: when looking at a single frame, you should be able to see at a glance what the algorithm just did. The caption is a parallel *textual* track for layouts that want narration apart from the animation.

This is why `last(frames)` defaults to `caption: false` — the inline note already labels the node, so adding a caption block below the static figure would be redundant. `stacked` and `figures` default to `caption: true` because animated rendering invites the textual narrative.

3.3. The six-frame rotation

`rotate-display` walks through rotation as six discrete frames: `init`, `pivots`, `break`, `restructure`, `connect`, `settle`. The sequence is calibrated for teaching: each frame answers exactly one question.

`step.kind`	Question answered
`init`	What's the starting tree?
`pivots`	Which two nodes are rotating?
`break`	Which edges are about to disappear?
`restructure`	What does the new shape look like, edges aside?
`connect`	Where do the new edges go?
`settle`	What does the final, clean tree look like?

An earlier draft had a "dashed red" intermediate between `pivots` and `break` (edges going dashed before disappearing). That was dropped because the orange pivot highlight already cues the eye to those edges — the extra frame added length without information.

The `restructure` frame is the load-bearing one: it shows the new tree shape with the rotated edges still hidden, so the student can see the new positions form *before* the edges connect. Without it, the rotation would collapse into one cut from "broken" to "rotated + green edges", which is too much information at once.

3.4. Path-based anchors and the cetz tree quirk

`cetz.tree.tree` numbers its nodes internally as `0`, `0-0`, `0-1`, `0-0-1`, … — sibling indices joined by hyphens, with the root being `0`. Starling configures cetz-tree's `group-name-prefix` to `"node-"`, so the resulting anchor names look like `node-0-0-1`. The whole tree also sits inside an outer named group (`"tree"` by default), so the fully qualified anchor for path `"LR"` is `tree.node-0-0-1`.

Two things made me reach for `path-anchor` rather than letting users type that out:

1. The user thinks in `"LR"`, not `"0-0-1"`. The translation is mechanical (`L → 0`, `R → 1`, prepend the root `0`) but it's error-prone to do by hand.
2. The naming scheme is a cetz-tree implementation detail. Wrapping it in `path-anchor(path, tree-name:, prefix:)` lets us swap to a different naming scheme later (e.g. if we drop cetz-tree for a custom layout) without churn at every annotation call site.

Phantom siblings (used to off-centre lone children — see `_build-cetz-tree`) also get cetz-tree-internal anchor names like `node-0-0-0`, but they are hidden at draw time and you generally should not annotate them.

4. BST animations tour

The BST class provides five `*-display` methods. Each returns `Array(Frame)`; the examples below pass the result to `starling.stacked` to render every frame vertically with its caption.

We'll use the same sample tree throughout this section:

```
#let t = (BST.new)(value: 4, label: auto, left: none, right: none)
#let t = (t.insert-many)(1, 7, 3, 6, 8)
```

4.1. Static display

`(t.display)()` returns a one-element frame array — the unmodified tree, with no step metadata.

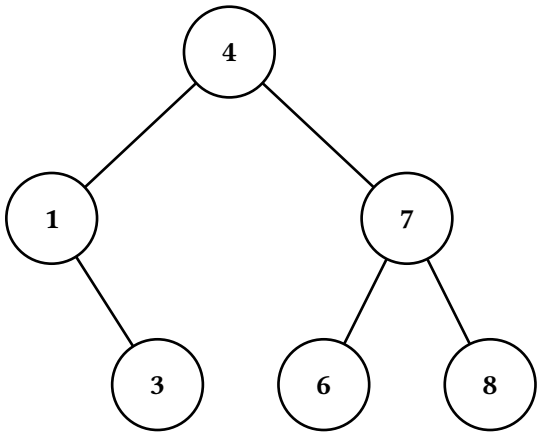

4.2. Search

`(t.search-display)(v)` walks the search path, highlighting each node visited and labelling it with the comparison made. `step.kind` is `"init"` for the initial frame and `"compare"` for each step; `step.found` records whether the value was reached.

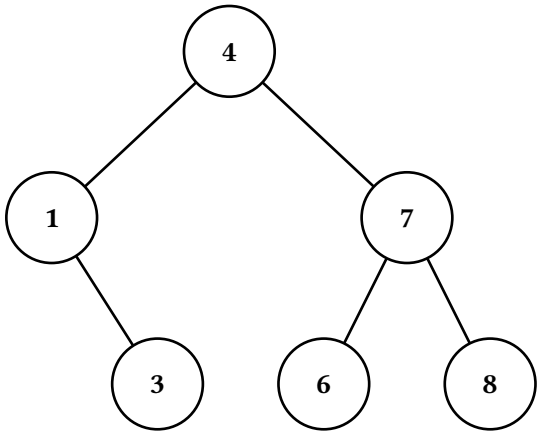

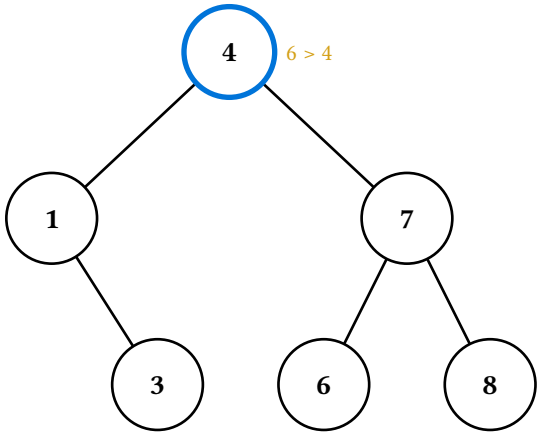

6 > 4

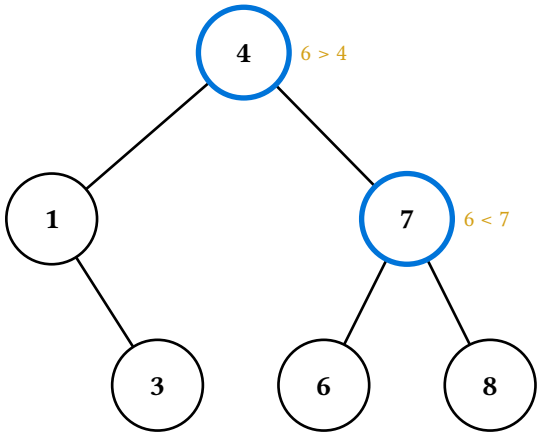

6 < 7

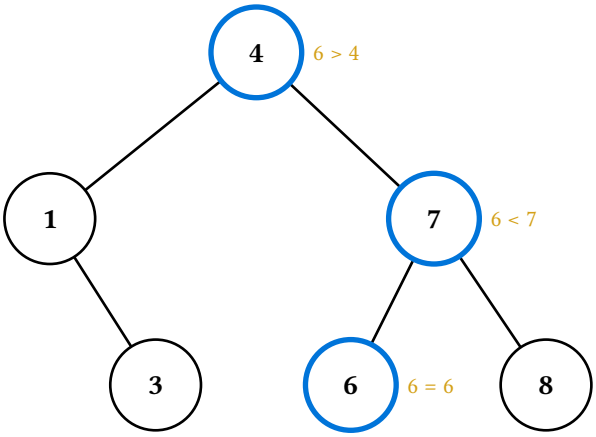

6 = 6

4.3. Insert

`(t.insert-display)(v)` walks the search path as for `search-display`, then transitions to the after-tree with the new node highlighted green. `step.kind` is `"init"`, `"compare"` (per search step), and finally `"inserted"`.

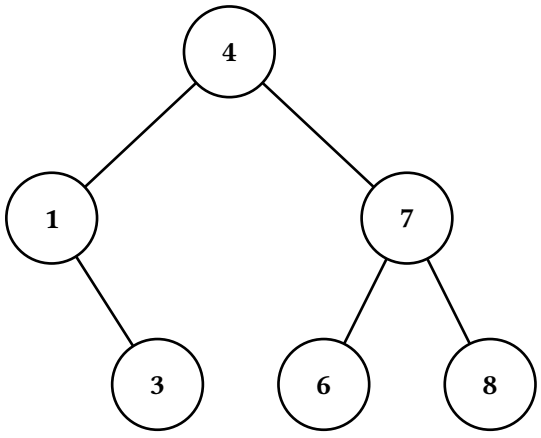

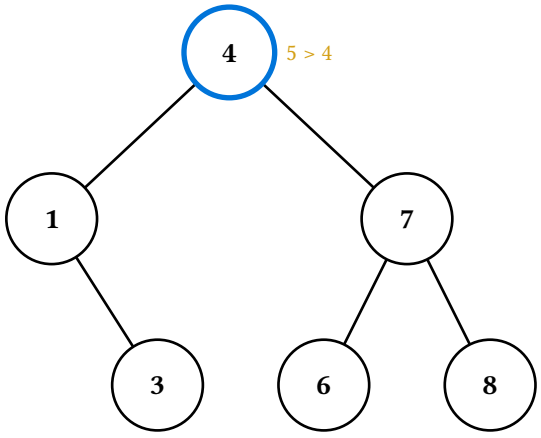

5 > 4

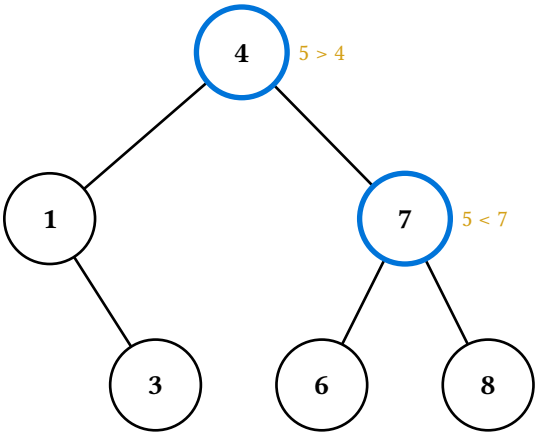

5 < 7

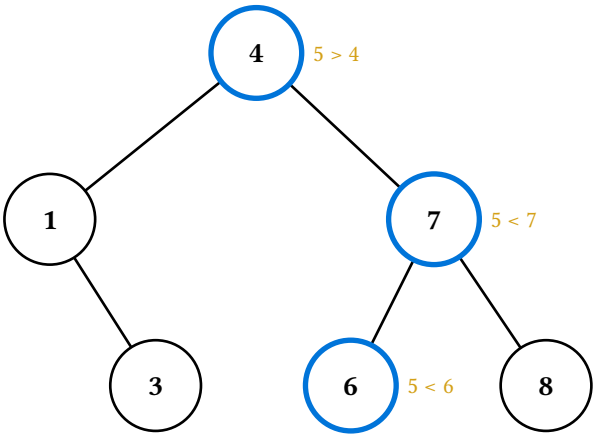

5 < 6

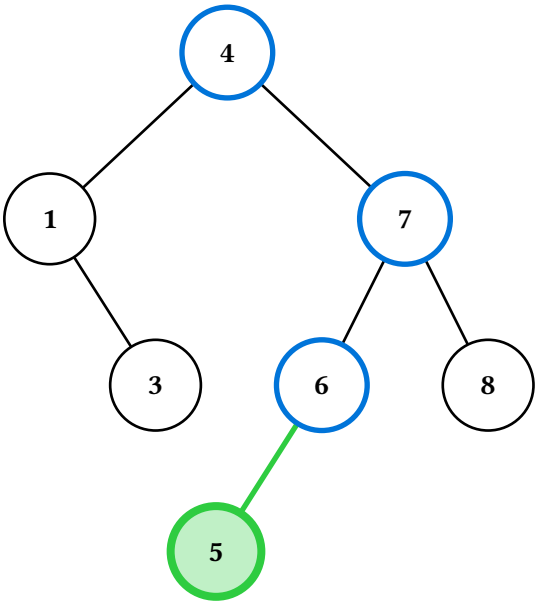

Inserted 5

4.4. Delete

`(t.delete-display)(v)` has three internal cases — leaf, one-child, and two-children — and dispatches on the target's children. `step.kind` ranges over `"init"`, `"highlight"`, `"break"`, `"descend"`, `"transfer"`, and `"settle"` depending on case.

The two-children deletion is the most complex: it highlights the target, descends into the left subtree to find the in-order predecessor, annotates the value transfer, then jumps to the after-tree with the new root of the affected subtree highlighted green.

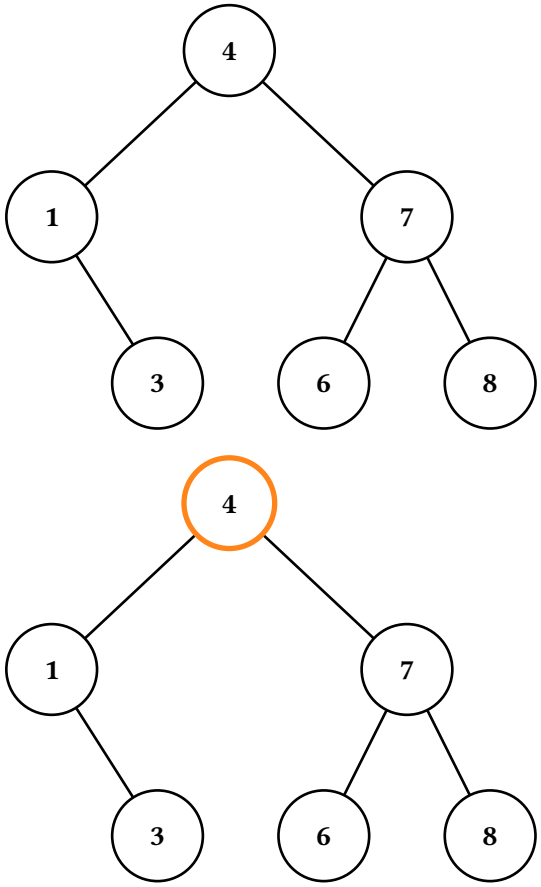

Delete 4

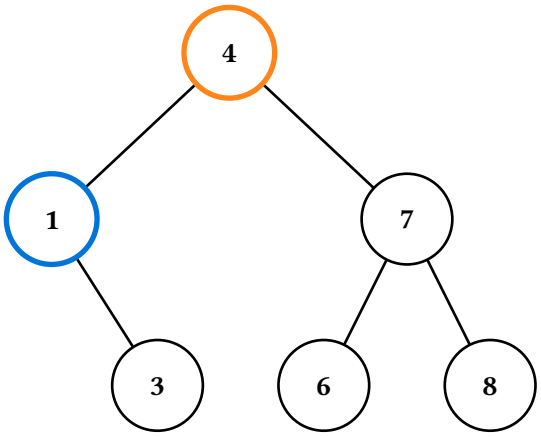

Find predecessor

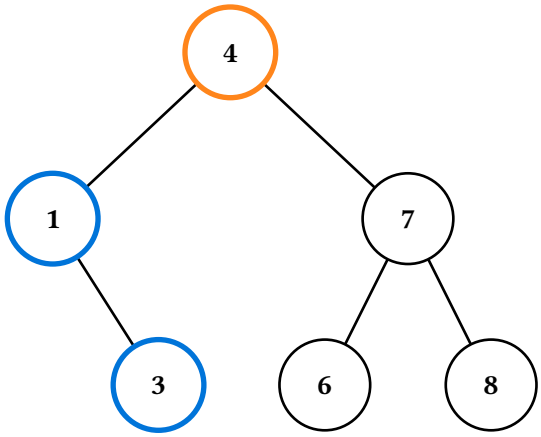

Find predecessor

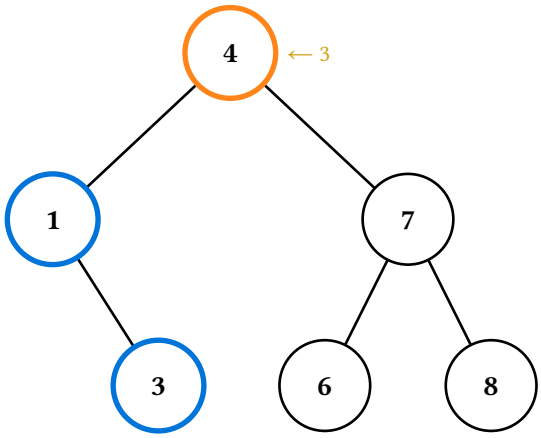

Transfer 3

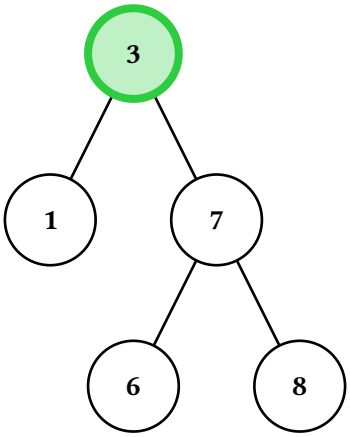

Done

4.5. Rotate

`(t.rotate-display)(c)` rotates around the direct child whose value is `c`. The six-frame sequence is described in Section 3.3.

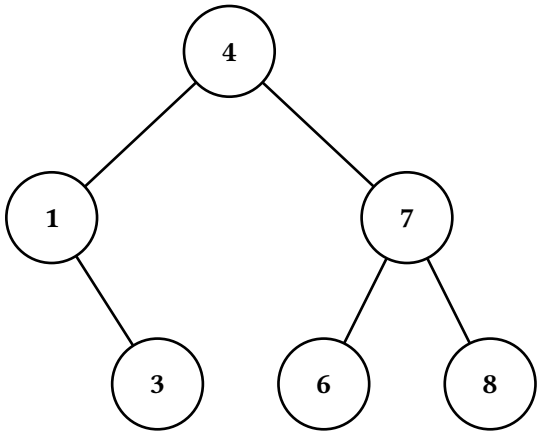

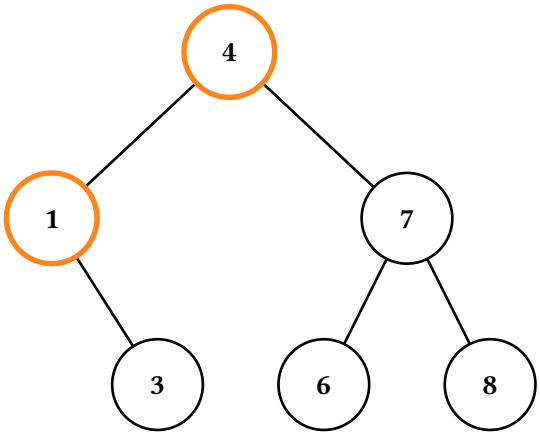

Rotate around 1

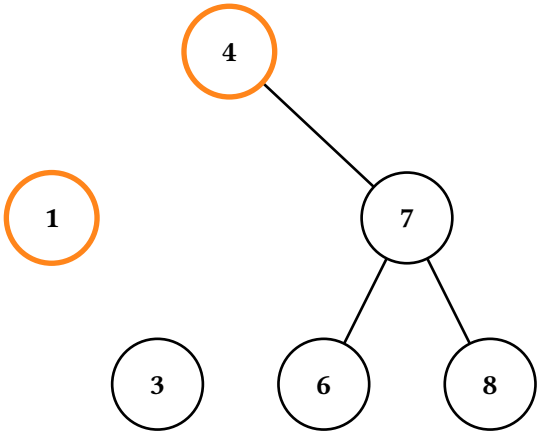

Break edges

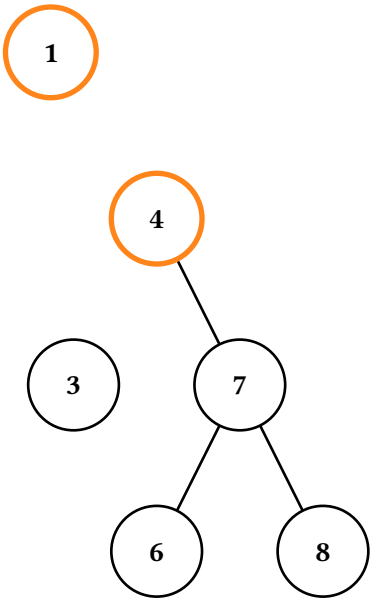

Restructure tree

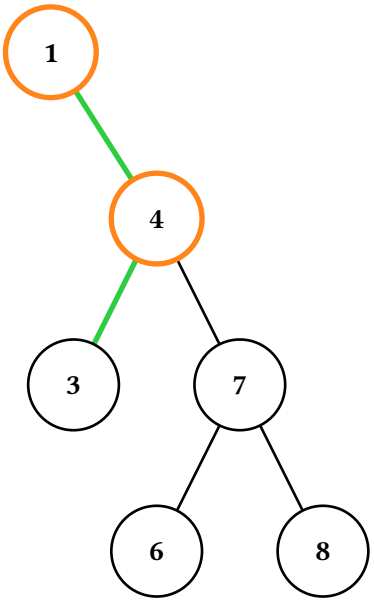

Reconnect edges

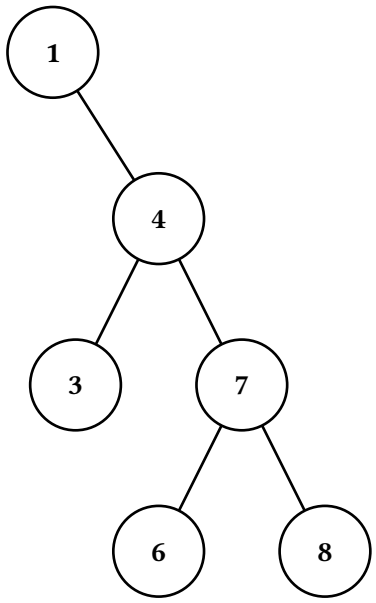

4.6. Traversals

`(t.in-order-display)()`, `(t.pre-order-display)()`, `(t.post-order-display)()`, and `(t.level-order-display)()` animate the four standard traversals. Each returns one frame per visit (plus the initial frame): the visited node is filled with the next color in a perceptually-uniform palette (magma by default) and tagged with a numbered badge marking its visit order, and the caption accumulates the running output sequence. `step.kind` is `"init"` for the initial frame and `"visit"` thereafter, with `step.value` and `step.index` (1-indexed) recording each visit. Pass a different `palette:` argument (e.g. `palette: color.map.viridis`) to switch color maps.

Pure-data variants `(t.in-order)()`, `(t.pre-order)()`, `(t.post-order)()`, and `(t.level-order)()` return the visit sequence as an array of path strings without producing any animation, in case you want to drive a custom layout.

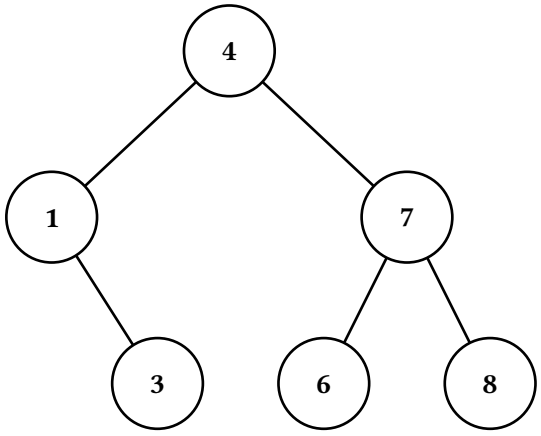

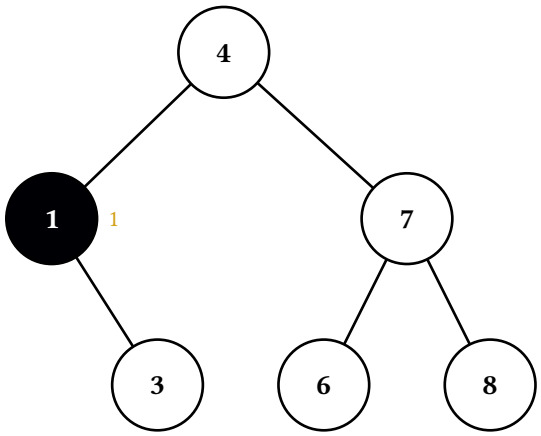

Output: `[1]`

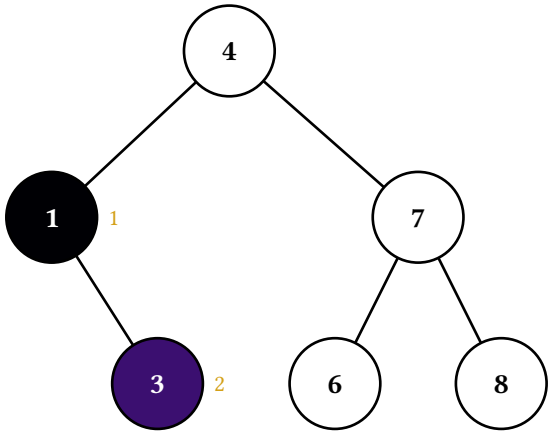

Output: `[1, 3]`

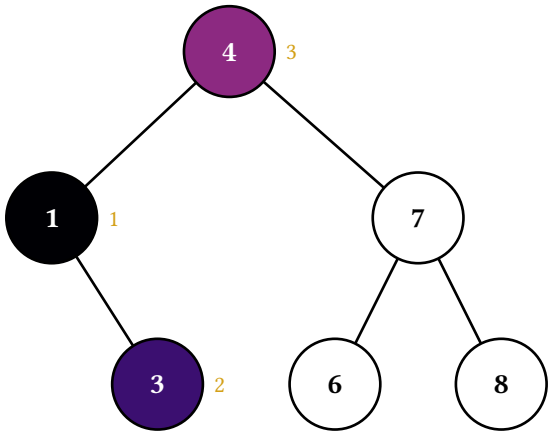

Output: `[1, 3, 4]`

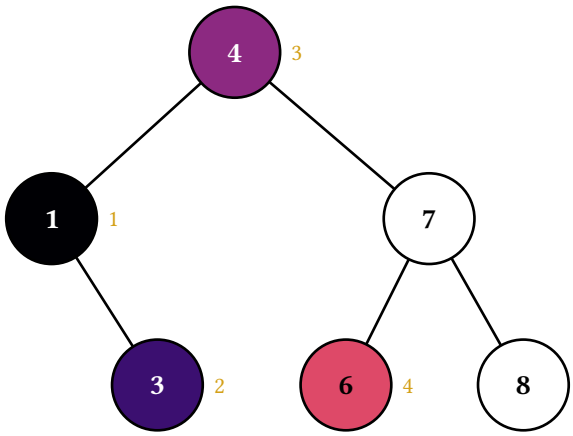

Output: `[1, 3, 4, 6]`

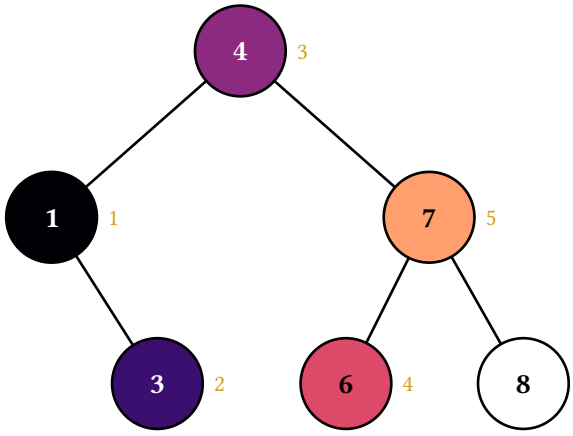

Output: `[1, 3, 4, 6, 7]`

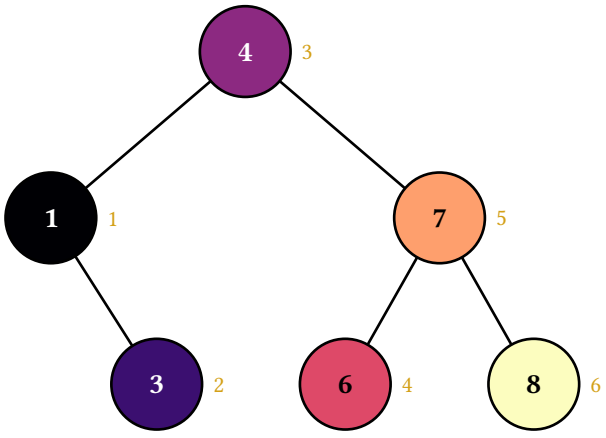

Output: `[1, 3, 4, 6, 7, 8]`

The final frame of each traversal carries the algorithm's full color signature. Laying all four out side-by-side gives a fast visual comparison — especially useful as a recall cue in review material where students already understand the algorithms and the gradient pattern reads as "oh right, post-order goes leaves-cool to root-warm":

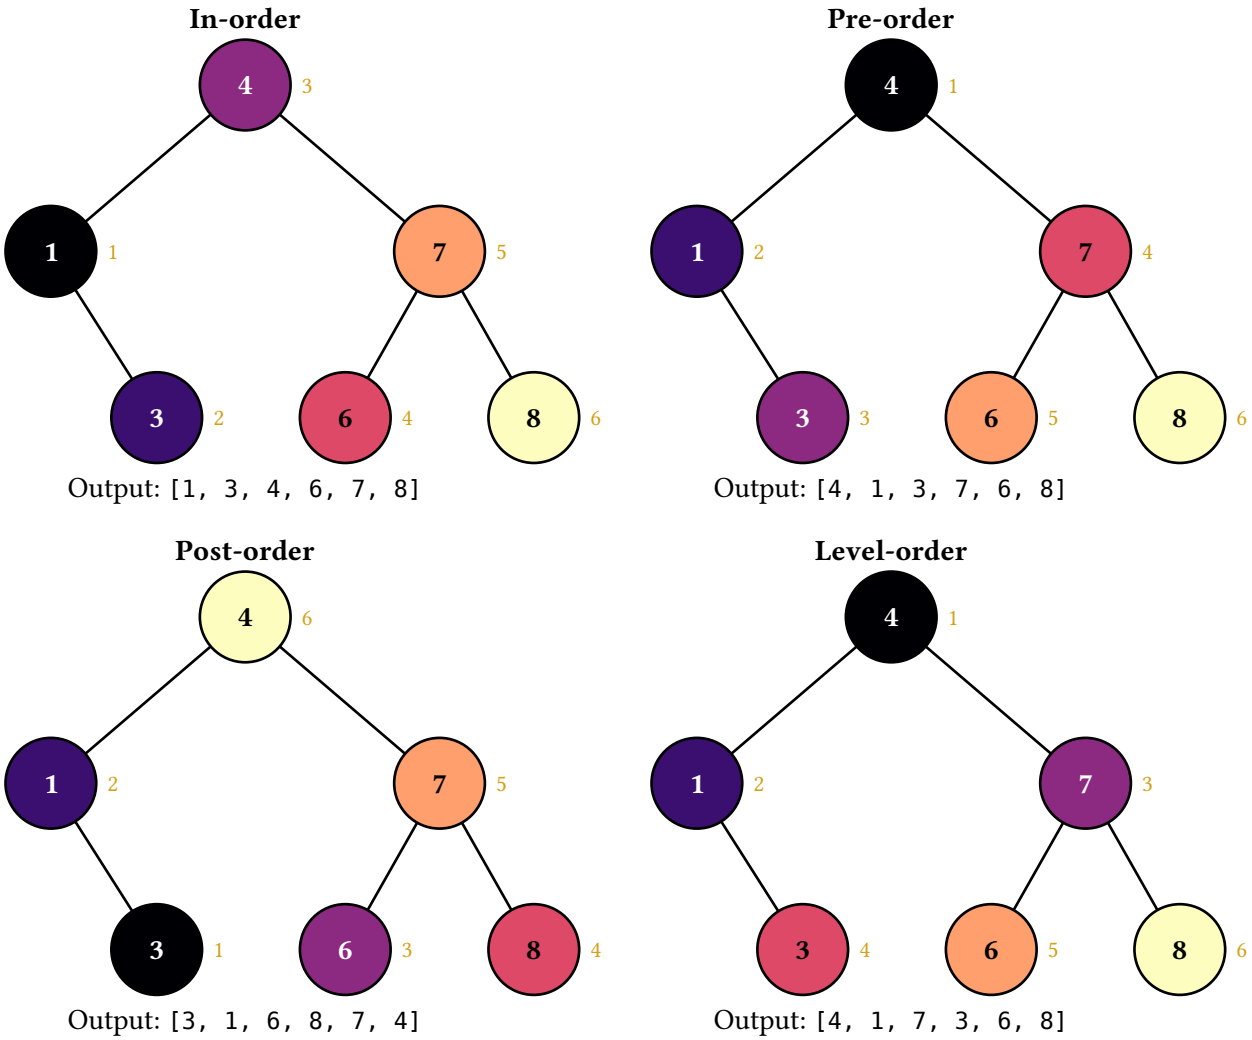

5. Touying composition examples

Starling is layout-agnostic. In a touying deck, the simplest use splats `figures` into `alternatives`:

```
#import "@preview/touying:0.7.3": *
#import "@preview/starling:0.2.0" as starling
#import starling: BST

== Searching
#alternatives(..starling.figures((t.search-display)(6)))
```

Each frame becomes a subslide; touying handles the rest.

5.1. Captions in a sidebar

When the visual animation should sit alongside synchronised text elsewhere on the slide, pull the captions out separately:

```
== Searching for 6
#let frames = (t.search-display)(6)
#grid(columns: 2, column-gutter: 2em,
  alternatives(..starling.figures(frames, caption: false)),
  alternatives(..frames.map(f => f.caption)),
)
```

Both `alternatives` calls produce the same number of subslides, so they step in lockstep.

5.2. Driving layout from step metadata

`frame.step` is free-form per-method metadata. For example, a slide that wants to colour-code its narration by step kind:

```
#let kind-colors = (compare: blue, inserted: green)

#let captioned-frames = frames.map(f => {
  let color = kind-colors.at(f.step.kind, default: black)
  figure(
    stack(dir: ttb, spacing: 0.5em,
      f.canvas, text(fill: color, f.caption)),
  )
})

#alternatives(..captioned-frames)
```

6. Composing with cetz annotations

For callouts or overlays that need to track specific nodes, drop down to the cetz layer. `draw-tree(tree, snapshot, ...)` emits the tree's cetz drawables without wrapping them in `cetz.canvas`, so you can compose them with your own draw commands in a shared canvas. `path-anchor(path)` translates an L/R path string to the cetz anchor name of the corresponding node.

```
#import "@preview/cetz:0.5.2"

#cetz.canvas({
  starling.draw-tree(t, starling.blank-snapshot())
  import cetz.draw: *
  circle(starling.path-anchor("LR"), radius: 0.85, stroke: red + 2pt)
})
```

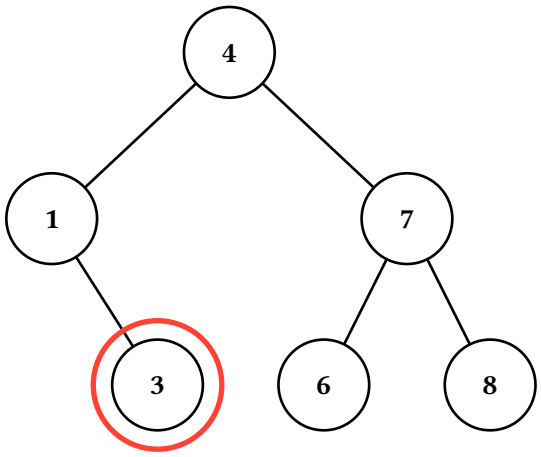

The same trick combines with frame canvases: each frame's `canvas` is itself a `cetz.canvas` block, so for static "annotate the final frame" use you can either render once via `draw-tree` (as above) or extract the cetz block from a frame and add to it externally.

7. The Op command stream

For animations that don't fit the BST's built-in methods — say, a custom hash-table probe sequence or a graph traversal — drop down to the `Op` command stream. Build a sequence of declarative ops and fold them into a renderer with `apply-ops`:

```
#let r = starling.make-renderer(t, sticky: true)
#let r = (r.with-caption)([start at root])
#let r = starling.apply-ops(r, (
  (starling.Op.Highlight.new)(path: "", color: blue),
  (starling.Op.Annotate.new)(path: "", text: [1 < 4]),
))
#let r = (r.push-frame)()
#let r = (r.with-caption)([descend left])
#let r = starling.apply-ops(r, (
  (starling.Op.ClearNotes.new)(),
  (starling.Op.Highlight.new)(path: "L", color: blue),
  (starling.Op.StyleEdge.new)(path: "L", style: (stroke: blue + 2pt)),
  (starling.Op.StyleNode.new)(path: "L", style: (fill: green.lighten(70%))),
  (starling.Op.Annotate.new)(path: "L", text: [1 = 1]),
))
#starling.stacked((r.render)())
```

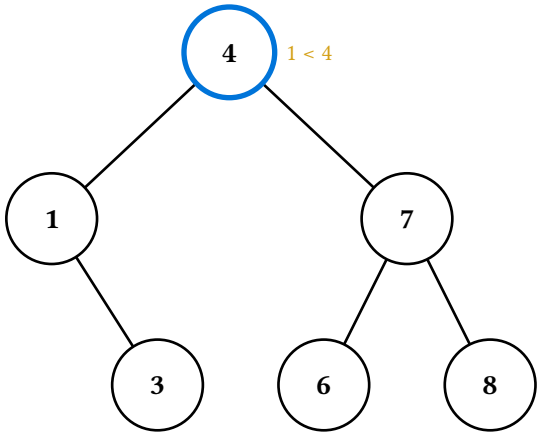

start at root

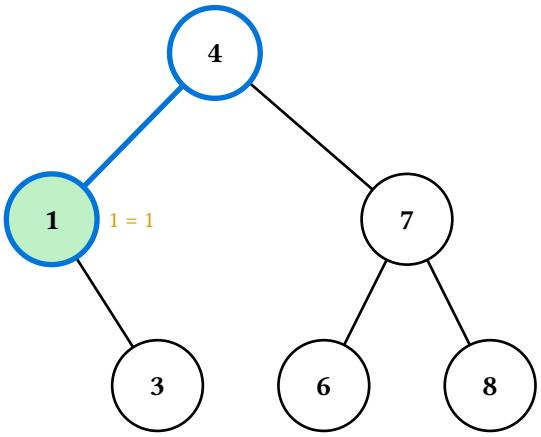

descend left

The trailing in-progress frame is kept implicitly — no final `Op.Commit` is needed.

`Op.Caption` does *not* exist; captions and step metadata are renderer-level, set via `r.with-caption(...)` / `r.with-step(...)` directly between Op batches. The split keeps each Op's responsibility narrow (modifying the current snapshot) and matches how the `*-display` methods are written internally.

8. API reference

The remainder of this manual is an auto-generated reference, produced by the `tidy` package from doc-comments in the source.

8.1. Render helpers

- canvases-only()
- figures()
- last()
- stacked()

Variables

- BST

8.1.1. canvases-only

Drop captions and step metadata, returning just the array of canvases. Useful when you want to lay out captions yourself — e.g. alongside other slide content in a custom touying layout.

8.1.1.1. Parameters

```
canvases-only(frames: array) -> array
```

8.1.1.1.1. frames `array`

The frame array returned by any `*-display` method.

8.1.2. figures

Build an array of `figure` content, one per frame, suitable for splatting into touying's `alternatives(..)` so each frame becomes its own subslide. Starling itself does not depend on touying — the result is just an array of standard Typst figures.

8.1.2.1. Parameters

```
figures(
  frames: array,
  caption: bool,
  alt: auto str,
  caption-spacing: length
) -> array
```

8.1.2.1.1. frames `array`

The frame array returned by any `*-display` method.

8.1.2.1.2. caption `bool`

Whether to stack each frame's caption below its canvas inside the figure body.

Default: `true`

8.1.2.1.3. alt `auto` or `str`

Alt-text override applied to every frame. `auto` uses each frame's `alt` field, falling back to the caption (if it's a plain string) and then to a generic `"Animation frame N"` placeholder.

Default: `auto`

8.1.2.1.4. caption-spacing `length`

Spacing between canvas and caption when `caption: true`.

Default: `0.5em`

8.1.3. last

Final frame as a single piece of content. Useful for static renderings in print or for "just show me the end state" usage.

The inline node annotations (drawn next to nodes in the cetz canvas) already label what the final step did, so the textual caption is suppressed by default. Pass `caption: true` to stack the caption below the canvas.

The returned content is wrapped in an alt-tagged figure for accessibility. The alt text comes from the frame's `alt` field; pass an explicit string to `alt` to override.

8.1.3.1. Parameters

```
last(
  frames: array,
  caption: bool,
  spacing: length,
  alt: auto str
) -> content
```

8.1.3.1.1. frames `array`

The frame array returned by any `*-display` method.

8.1.3.1.2. caption `bool`

Whether to render the step's caption below the canvas.

Default: `false`

8.1.3.1.3. spacing `length`

Vertical spacing between canvas and caption when `caption: true`.

Default: `0.5em`

8.1.3.1.4. alt `auto` or `str`

Alt-text override. `auto` uses the frame's `alt` field.

Default: `auto`

8.1.4. stacked

All frames stacked vertically as one block of content. Useful for handouts that want to show the full animation in a single figure. Captions are on by default since the stack is read as a sequence — the caption text annotates which step each tree corresponds to.

Each frame is wrapped individually in an alt-tagged figure so screen-reader users hear the per-step narration. Pass a string to `alt` to set the same override on every frame.

8.1.4.1. Parameters

```
stacked(
  frames: array,
  caption: bool,
  spacing: length,
  caption-spacing: length,
  alt: auto str
) -> content
```

8.1.4.1.1. frames `array`

The frame array returned by any `*-display` method.

8.1.4.1.2. caption `bool`

Whether to include each frame's caption below its canvas.

Default: `true`

8.1.4.1.3. spacing `length`

Vertical spacing between frames.

Default: `1em`

8.1.4.1.4. caption-spacing `length`

Spacing between each canvas and its caption.

Default: `0.5em`

8.1.4.1.5. alt `auto` or `str`

Alt-text override applied to every frame. `auto` uses each frame's `alt` field.

Default: `auto`

8.1.5. BST

The default BST class. Build a tree via `(BST.new)(value:, label:, left:, right:)` — set `label: auto` to render the integer `value` as the drawn label, or pass content (a string, image, or arbitrary content) for custom labels. Call its `*-display` methods (see the BST API section of the manual) to get animation frames; pass those frames through one of the helpers below.

8.2. Animation kernel

- apply-ops()
- blank-snapshot()
- concat-frames()
- draw-tree()
- make-renderer()
- path-anchor()

Variables

- PathId
- NodeStyle
- EdgeStyle
- Snapshot
- Frame
- TreeRenderer
- Op

8.2.1. apply-ops

Fold a sequence of `Op` values into a renderer, returning the updated renderer. The trailing in-progress frame is kept — no explicit `Op.Commit` is needed after the last batch of edits.

8.2.1.1. Parameters

```
apply-ops(
  renderer: TreeRenderer,
  ops: array
) -> TreeRenderer
```

8.2.1.1.1. renderer `TreeRenderer`

The renderer to apply ops to.

8.2.1.1.2. ops `array`

Sequence of `Op` values to fold left-to-right.

8.2.2. blank-snapshot

Build a fresh `Snapshot` with no node or edge style overrides. Useful when you want to render a tree once with `draw-tree` and no per-frame state.

8.2.2.1. Parameters

```
blank-snapshot() -> Snapshot
```

8.2.3. concat-frames

Stitch frame arrays from multiple renderers (or pre-rendered arrays) into one flat `Array(Frame)`. Use this when an animation spans a tree-shape change (rotation, deletion, insertion) and one renderer can't hold all the frames — one renderer per shape, joined at the boundary.

Accepts a mix of `TreeRenderer` instances and already-rendered arrays; the former are `render`-ed in place.

8.2.3.1. Parameters

```
concat-frames(..parts) -> array
```

8.2.3.1.1. ..parts

Variadic mix of renderers and frame arrays.

8.2.4. draw-tree

Emit the cetz draw commands for one styled snapshot of `tree`, *without* wrapping them in a `cetz.canvas`. Caller is responsible for the canvas wrap. Use this when you want to add your own cetz annotations alongside the tree — for example, callouts anchored at specific nodes.

The whole tree is wrapped in a named cetz group (`name`, default `"tree"`); each non-phantom node has an inner-name of `<node-prefix><cetz-tree-name>` where the cetz-tree name is `"0"` for the root, `"0-0"` for L, `"0-1"` for R, `"0-0-1"` for LR, and so on. Fully qualified anchor names therefore look like `"tree.node-0-0-1"` — use `path-anchor()` to compute them from an L/R path string.

8.2.4.1. Parameters

```
draw-tree(
  tree: dictionary,
  snapshot: Snapshot,
  default-node-style: dictionary,
  default-edge-style: dictionary,
  name: str,
  node-prefix: str
) -> content
```

8.2.4.1.1. tree `dictionary`

The tree to render — any value with `value`, `label`, `left`, and `right` fields (e.g. a `BST` instance). `label` of `auto` falls back to `str(value)` any other value (string, image, content) is rendered as the node's drawn label.

8.2.4.1.2. snapshot `Snapshot`

Style overlay for this snapshot. Use `blank-snapshot()` for an unstyled tree.

8.2.4.1.3. default-node-style `dictionary`

Default node-style overrides applied before per-node overrides.

Default: `(:)`

8.2.4.1.4. default-edge-style `dictionary`

Default edge-style overrides applied before per-edge overrides.

Default: `(:)`

8.2.4.1.5. name `str`

Outer group name for the whole tree. Must match the `tree-name` argument to `path-anchor()` for anchors to resolve.

Default: `"tree"`

8.2.4.1.6. node-prefix `str`

Prefix attached to each node's cetz-tree internal name. Must match the `prefix` argument to `path-anchor()`.

Default: `"node-"`

8.2.5. make-renderer

Build a `TreeRenderer` seeded with one blank initial frame.

With `sticky: true` (the default), each new frame pushed via `r.push-frame()` starts from the previous frame's style snapshot — so highlights accumulate over the course of an animation. Set `sticky: false` to start each frame from a clean slate.

8.2.5.1. Parameters

```
make-renderer(
  tree: dictionary,
  default-node-style: dictionary,
  default-edge-style: dictionary,
  sticky: bool
) -> TreeRenderer
```

8.2.5.1.1. tree `dictionary`

The tree to render — any value with `value`, `label`, `left`, and `right` fields. See `draw-tree()` for how `label` is rendered.

8.2.5.1.2. default-node-style `dictionary`

Default node-style overrides applied before per-snapshot overrides.

Default: `(:)`

8.2.5.1.3. default-edge-style `dictionary`

Default edge-style overrides applied before per-snapshot overrides.

Default: `(:)`

8.2.5.1.4. sticky `bool`

When `true`, new frames inherit the previous frame's style. When `false`, each new frame starts blank.

Default: `true`

8.2.6. path-anchor

Translates a starling path-id (L/R string rooted at `""`) to the fully-qualified cetz anchor name produced by `draw-tree()`. Path `""` maps to the root anchor; `"L"`/`"R"` map to first-level children; and so on.

The `tree-name` and `prefix` arguments must match the `name` and `node-prefix` passed to `draw-tree`.

8.2.6.1. Parameters

```
path-anchor(
  path: str,
  tree-name: str,
  prefix: str
) -> str
```

8.2.6.1.1. path `str`

L/R path string identifying a node (`""` = root, `"LR"` = root's left child's right child, etc.).

8.2.6.1.2. tree-name `str`

Outer-group name; must match `draw-tree`'s `name`.

Default: `"tree"`

8.2.6.1.3. prefix `str`

Per-node prefix; must match `draw-tree`'s `node-prefix`.

Default: `"node-"`

8.2.7. PathId

Typsy refinement: a string built only from `"L"` and `"R"` characters. The empty string is the root; `"L"` is the root's left child; `"LR"` is the root's left child's right child; and so on.

8.2.8. NodeStyle

Typsy refinement: a dictionary of node-style overrides. Recognised keys are `fill`, `stroke`, `text-fill`, `note`, `note-fill`, and `hide`.

8.2.9. EdgeStyle

Typsy refinement: a dictionary of edge-style overrides. Recognised keys are `stroke`, `note`, `note-fill`, `mark`, and `hide`.

8.2.10. Snapshot

Typsy class representing one sparse style overlay — the per-node and per-edge style overrides for a single frame. Build one via `blank-snapshot()` and chain its `style-node` / `style-edge` / `note-node` / `note-edge` / `clear-notes` methods. Snapshots are rarely constructed directly — they're managed by `make-renderer()` and the BST's `*-display` methods.

8.2.11. Frame

Typsy class representing one rendered animation frame. Fields: `canvas` (cetz content, no caption baked in), `caption` (optional textual track for the step, possibly `none`), `step` (optional free-form per-method metadata, possibly `none`), and `alt` (optional accessible text describing what the frame depicts, possibly `none`). Produced by `TreeRenderer.render()` the BST's `*-display` methods return arrays of these.

The `alt` field carries text destined for the `alt:` argument on a Typst `figure`. The first frame of each operation includes the full tree structure (via the BST `describe` method) so screen-reader users have a baseline; subsequent frames describe only the per-step change.

8.2.12. TreeRenderer

Typsy class accumulating an animation. Holds the tree, parallel arrays of snapshots / captions / steps / alts (one entry per frame), and default node/edge styles. Methods include `push-frame()`, `patch(fn)`, `push-with-node(path, ..style)`, `push-with-edge(path, ..style)`, `push-note-node(path, txt)`, `push-note-edge(path, txt)`, `with-caption(c)`, `with-step(s)`, `with-alt(a)`, and `render()`. Build one via `make-renderer()`.

8.2.13. Op

Typsy enumeration describing a declarative command stream for building animations. Variants:

- `Op.Highlight(path, color)` — set a node's stroke to `color + 2pt`.
- `Op.Annotate(path, text)` — attach an inline note to a node (drawn in-canvas).
- `Op.StyleNode(path, style)` — arbitrary node-style overrides.
- `Op.StyleEdge(path, style)` — arbitrary edge-style overrides.
- `Op.Commit()` — finalise the current frame and open a fresh one.
- `Op.ClearNotes()` — drop all notes from the current frame's snapshot.

Captions and step metadata are renderer-level rather than snapshot-level: set them between Op batches via `r.with-caption(...)` / `r.with-step(...)`, not through an Op. Fold a sequence of Ops into a renderer with `apply-ops()`.